

Exponential-Range Mass Production of Boolean Functions

Local recovery, disjoint scheduling, and bounded congestion

Samuel Schlesinger

September 2026

Abstract

We study the cost of evaluating one Boolean function on many independent inputs. For $f : \{0, 1\}^n \rightarrow \{0, 1\}$, let $f^{\times t}$ evaluate f on t such inputs. Uhlig showed that $t = 2^{o(n/\log n)}$ copies preserve the sharp one-copy asymptotic. We retain the optimal order of growth throughout every fixed exponential range below the 2^n -copy scale: for every fixed $0 \leq \gamma < 1$,

$$C(f^{\times t}) = O_\gamma(2^n/n) \quad \text{for all } t \leq 2^{\gamma n}.$$

The bound is uniform in f and t . It does not preserve Uhlig’s leading constant: in the standard basis used here, a high-rate version of our construction gives coefficient at most $1/(1 - \gamma)$, while the sharp one-copy coefficient is 1.

The proof uses local coded recovery instead of scanning the whole truth table. An evaluation code gives each requested symbol many affine-line recovery sets. A deterministic scheduler chooses disjoint lines within request groups, so each shorter resource function is queried at most once per group; induction handles the demand across groups. With $k + 1$ equal input blocks this gives the copy exponents

$$\frac{1}{2}, \frac{2}{3}, \frac{3}{4}, \dots,$$

which approach one.

This work was developed with substantive AI assistance. It began in a ChatGPT.com conversation that the author exported as a $\text{\LaTeX}$ document. Later interactive sessions used Claude Fable 5 and 5.1 through Claude Code, and ChatGPT 5.6 Sol through Codex. These systems contributed to experiments, broad literature searches, and several core parts of the proof strategy and its internal constructions. The author checked the arguments and citations and is responsible for the paper.

1 Introduction

How much circuit size can be saved when the same function must be evaluated on many unrelated inputs? This is a direct-sum question in which the outputs are independent but the circuit implementing them need not be: intermediate computations may be reorganized and shared across copies.

For a Boolean function $f : \{0, 1\}^n \rightarrow \{0, 1\}$ and an integer $t \geq 1$, define

$$f^{\times t}(x^{(1)}, \dots, x^{(t)}) = (f(x^{(1)}), \dots, f(x^{(t)})).$$

The naive construction uses t disjoint copies of a circuit for f and gives

$$C(f^{\times t}) \leq tC(f).$$

The mass-production problem asks how large t can become before the cost must rise above the worst-case scale for one copy. For unrestricted circuits, the naive direct-sum bound can be far from optimal. Uhlig's theorem allows $t = 2^{o(n/\log n)}$ independent evaluations with no asymptotic increase in the leading worst-case one-copy synthesis cost. In our basis and cost convention that coefficient is 1 [5]; its numerical value can depend on the basis and convention. See Uhlig's later account [14], Wegener's Theorem 10.2.3 [15], and the modern restatement in [8]; the 1974 paper [13] is a precursor.

We prove, by a different construction, that the order-of-growth plateau reaches $2^{\gamma n}$ simultaneous evaluations for every fixed $\gamma < 1$. The price is that we control only the optimal worst-case order $2^n/n$, not the sharp one-copy leading constant.

1.1 Main result

We fix the standard basis $\{\wedge, \vee, \neg\}$, with fan-in two for the binary gates. Fan-out and output designations are free, and circuit size is the number of non-input gates. Any other fixed finite complete basis changes the result by at most a constant-factor simulation.

Theorem 1.1 (Exponential-range mass production). *For every fixed $0 \leq \gamma < 1$, there is a constant A_γ such that, for every integer $n \geq 1$, every Boolean function $f : \{0, 1\}^n \rightarrow \{0, 1\}$, and every integer $1 \leq t \leq 2^{\gamma n}$,*

$$C(f^{\times t}) \leq A_\gamma \frac{2^n}{n}.$$

The coefficient hidden in this statement can also be controlled. The high-rate variant in [Proposition A.3](#) proves

$$\limsup_{n \rightarrow \infty} \max_{\substack{f: \{0,1\}^n \rightarrow \{0,1\} \\ 1 \leq t \leq 2^{\gamma n}}} \frac{nC(f^{\times t})}{2^n} \leq \frac{1}{1 - \gamma}. \quad (1)$$

The right side is not claimed to be sharp. Our conjecture is that it can be replaced by the one-copy coefficient 1 for every fixed $\gamma < 1$.

The order of the quantifiers matters. Once γ is fixed, one constant A_γ works for every function and every batch size in the stated range; the bound has no factor depending on t . At the largest allowed batch size, the amortized size per requested value is therefore

$$\frac{C(f^{\times t})}{t} = O_\gamma \left(\frac{2^{(1-\gamma)n}}{n} \right).$$

The theorem does not make A_γ uniform as $\gamma \rightarrow 1$.

There must nevertheless be a transition before the 2^n -copy scale. If $n \geq 2$ and f depends on all n variables, every circuit for $f^{\times t}$ has size at least $tn/2$. Hence a uniform $O(2^n/n)$ plateau cannot persist once $t = \omega(2^n/n^2)$; see [Section 8](#). Our theorem reaches every fixed exponential rate, but leaves this near-endpoint regime open.

1.2 How the construction works

Paul’s general many-copy construction provides a useful contrast [10, Lemma 2]. It treats the truth table of f as an ordered database: sort the t input addresses, sweep once through the hardwired table, and restore the answers to their original order. This global lookup strategy works for arbitrary t and gives

$$O\left(\max\{n2^n, nt \log^2 t\}\right).$$

For every fixed $\gamma < 1$ and $t \leq 2^{\gamma n}$, the full-table term $n2^n$ eventually dominates.

We instead turn the truth table into shorter resource functions and recover a requested value from a selectable local set of them. The quantity that has to be controlled is *simultaneous demand*: how many different suffixes ask for the same resource? Disjoint recovery sets control demand within a group, and recursion controls it across groups. Both Paul’s construction and ours use sorting, but for different jobs. Paul sorts input addresses before a global table scan; we sort scheduling and routing records after choosing the local recovery sets.

Here is the construction in the order in which a requested value passes through it.

1. **Pack restrictions.** Split each input as $x = (a, z)$. For every fixed suffix z , pack the values $f(a, z)$ over all prefixes a into the information symbols of a constant-rate low-degree evaluation code over $\mathbb{F}_q$.
2. **Recover locally.** A requested information symbol can be reconstructed from the $q - 1$ non-target points on any affine line through it. About $q^{\ell-1}$ directions are available, so every request has many alternative recovery sets.
3. **Schedule bounded congestion.** Partition the requests into C groups. A deterministic greedy circuit chooses pairwise disjoint recovery lines inside each group. Consequently, each code coordinate is requested on at most one suffix per group and at most C suffixes overall.
4. **Recurse on the suffix.** For a fixed code coordinate, evaluating its resource function on those at most C suffixes is itself a C -copy problem, now on only $d = |z|$ variables. The recursive hypothesis pays for this remaining simultaneous demand.

Section 3 first shows the disjoint-recovery idea in Uhlig’s two-copy construction. The local gadget is stated in Proposition 4.2, and Proposition 6.1 accounts for its circuit cost. The parameter choice is then isolated in Section 7. Gate-level details for direction selection and routing are deferred to Section C; the leading coefficient is tracked in Section A.

1.3 Why the exponents approach one

For an integer $k \geq 1$, let P_k be the claim that every fixed copy exponent below $k/(k + 1)$ has worst-case cost $O(2^n/n)$. To prove P_k , write

$$n = (k + 1)m + O_k(1)$$

and view the input as $k + 1$ blocks of length about m . The outer construction encodes one block for local recovery; the remaining k blocks support the recursive resource functions.

Write $\gamma = k/(k+1) - \eta$. In the new range not already covered by P_{k-1} , choose the number of request groups by

$$\log_2 C = (k-1)m - \frac{\eta n}{2} + O(1).$$

Since $\log_2 t \leq km - \eta n + O_k(1)$, each group contains at most

$$2^{m-\eta n/2+O_k(1)}$$

requests: strictly fewer than the one-block scale 2^m . At the same time, C has exponent strictly below $(k-1)/k$ when measured against the remaining k blocks, so P_{k-1} evaluates every resource. Fixed finite-field slack makes all bookkeeping terms exponentially smaller than 2^n .

The base case P_1 uses two equal blocks and has no recursive congestion. At every later level, one block supplies the scheduling room inside each group and the preceding statement is invoked on the other blocks.

1.4 Relation to earlier work

The closest asymptotic guarantees are summarized below. Here “one-copy leading term” means the basis-dependent Shannon–Lupanov worst-case asymptotic.

Result	Copy range	Circuit-size bound	Comparison
Uhlig [14, 15, 8]	$\log t = o(n/\log n)$	one-copy leading term	optimal order and leading constant in a smaller copy range
Paul, Lemma 2 [10]	arbitrary t	$O(\max\{n2^n, nt \log^2 t\})$	$O(n2^n)$ throughout every fixed exponential range below one
Albrecht, Theorem 1 [1]	$t \leq 2^n$	$O(2^n n^2)$	the full exponential range, with polynomial overhead
This paper	$t \leq 2^{\gamma n}$ for fixed $\gamma < 1$	$O_\gamma(2^n/n)$	optimal worst-case order, with a nonsharp γ -dependent coefficient

For fixed $\gamma < 1$ and $t \leq 2^{\gamma n}$, Paul’s second term is $o(n2^n)$, so his displayed bound is $O(n2^n)$. The theorem proved here improves that bound by a factor of n^2 at the level of polynomial prefactors. Paul’s g^k is the same k -fold product task considered here; his circuits permit arbitrary binary Boolean gates, which changes our fixed basis cost by at most a constant factor.

Albrecht studies the same simultaneous-realization problem over a complete binary basis [1, Theorem 1]. His construction handles every $t \leq 2^n$ with size $O(2^n n^2)$ and depth $O(n \log n)$. Since the one-copy Shannon size is $\Theta(2^n/n)$, this is an $O(n^3)$ multiplicative overhead, as Albrecht also notes. Thus his theorem covers the endpoint copy range but does not give a constant-factor plateau at the one-copy scale.

The coding interpretation begins with Uhlig’s two-copy construction, which is the length- $(2^p + 1)$ single-parity-check code viewed as a two-request disjoint recovery code. Systematic multivariate polynomial evaluation, recovery on affine lines, and the use of disjoint recovery sets appear in the batch-code literature, especially Section 3.4 of Ishai–Kushilevitz–Ostrovsky–Sahai [7]. Their geometric procedure chooses lines randomly and deletes intersections. The construction here instead gives a deterministic greedy selection circuit, partitions requests into groups, recursively absorbs the resulting inter-group congestion, and charges the scheduler and fixed-wire routing explicitly. We make no priority claim for these refinements.

Later lifted-polynomial constructions also use line-based recovery to obtain batch-code guarantees [6]. Private-information-retrieval codes of Fazeli, Vardy, and Yaakobi form a related family in which each information symbol has many disjoint recovery options [4]. The self-contained proof of Theorem 1.1 uses only the elementary product-code and affine-line facts proved below. The sharper coefficient bound (1) also uses the rate estimate for lifted Reed–Solomon codes from [6].

Independently of historical priority, the technical contribution established here is a self-contained deterministic disjoint-line scheduler, an explicit fixed-wire scatter/gather implementation, and the bounded-congestion composition that drives the equal-block induction.

1.5 AI-assisted development and provenance

The project began in an interactive conversation on ChatGPT.com, exported to the author’s computer as a $\text{\LaTeX}$ document. Later work used Claude Fable 5 and 5.1 through Claude Code and ChatGPT 5.6 Sol through Codex in a series of interactive sessions. These systems assisted with mathematical experiments, broad literature searches, and parts of the proof strategy and its internal constructions; their role went beyond prose editing. The author chose which directions to pursue, checked and revised the arguments and literature claims, and takes responsibility for the paper.

2 Circuit conventions and worst-case notation

Our circuits are finite directed acyclic graphs. Source vertices are input bits or the constants zero and one. Internal vertices are labelled by $\wedge$, $\vee$, or $\neg$, with the usual fan-in. A circuit for a map $F : \{0, 1\}^N \rightarrow \{0, 1\}^m$ has m designated output wires, which may be input wires or gate outputs and may repeat. Fan-out, constants, and output designations carry no cost; size is the number of internal gates. We write $C(F)$ for the minimum size of such a circuit. All leading-constant comparisons in this paper refer to this fixed model.

Let

$$L(n) = \max_{f: \{0,1\}^n \rightarrow \{0,1\}} C(f).$$

The Shannon counting lower bound and Lupanov’s synthesis upper bound imply

$$L(n) = \Theta(2^n/n)$$

for every fixed finite complete basis; see [12, 9, 15]. We use only the upper bound.

For $r \geq 1$, define

$$L_r(n) = \max_{f: \{0,1\}^n \rightarrow \{0,1\}} C(f^{\times r}).$$

For every $r \geq 1$, one has $L_r(n) \geq L(n)$: fix all but one of the r input blocks and retain the corresponding output. Thus $2^n/n$ is the optimal worst-case order even when several outputs are requested.

For each f , choose a circuit for $f^{\times r}$ of size at most $L_r(n)$. If a construction supplies fewer than r live inputs, the unused inputs of this circuit may be fixed arbitrarily and their outputs ignored. Thus a bound for $L_r(n)$ also applies to every smaller number of copies.

All circuits below are nonuniform. Field representations, irreducible polynomials, sorting networks, and other finite combinatorial objects may be fixed as part of the circuit for each input length.

We use $\tilde{O}_\ell(X)$ for

$$X \cdot \text{poly}_\ell(n, \log q, \log(t+2), \log(C+2)).$$

Only polynomial factors are suppressed; in the parameter regimes used below, they never affect an exponential rate.

For a positive integer k , write $[k]_0 = \{0, \dots, k-1\}$.

We repeatedly use two standard circuit primitives. First, after representing $\mathbb{F}_{2^b}$ in a polynomial basis, field addition, multiplication, inversion of a nonzero element, and comparison of canonical b -bit representations have circuits of size $\text{poly}(b)$. For example, inversion can be implemented by binary exponentiation to the power $2^b - 2$. Second, R records of width w can be sorted by a fixed sorting network using $O(R \log^2 R)$ compare-exchange units, each of size $O(w)$; see [3]. Sorting by a leading type bit, with original positions used as tie breakers when needed, also gives stable fixed-wire compaction within the same asymptotic bound.

3 The two-copy construction as disjoint recovery

Uhlig's two-copy argument already contains the combinatorial invariant used later. We will not use its numerical bound. What matters is that each requested value has more than one representation, allowing the two requests to choose representations that do not reuse a resource at the same time.

Fix a prefix length p , put $M = 2^p$, and define the restrictions

$$f_i(z) = f(i, z), \quad i \in \{0, \dots, M-1\}.$$

Uhlig's two-copy construction introduces the following functions; see the proof of Wegener's Theorem 10.2.3 [15] and the overview in [8, Section 1.2]:

$$g_0 = f_0, \quad g_i = f_{i-1} \oplus f_i \quad (1 \leq i < M), \quad g_M = f_{M-1}.$$

Every restriction has two representations:

$$f_i = \bigoplus_{j=0}^i g_j = \bigoplus_{j=i+1}^M g_j.$$

Given requests $i \leq j$, recover f_i from the prefix set $\{0, \dots, i\}$ and f_j from the suffix set $\{j+1, \dots, M\}$. The sets are disjoint, so each resource function g_j is evaluated on at most one suffix.

For two requests arriving in the opposite order, a comparator first orders their prefixes, conditional swaps route the two suffixes to the prefix and suffix recovery branches, and a final swap restores the output order. This costs only $\text{poly}(n)$ gates.

The map $(f_0, \dots, f_{M-1}) \mapsto (g_0, \dots, g_M)$ is, up to an invertible change of information coordinates, the length- $(M+1)$ single-parity-check code. The two representations above are disjoint recovery sets for each information coordinate. This suggests replacing the one-dimensional telescoping code by a constant-rate code with many geometric recovery sets.

The full construction keeps two features of this example. The resource functions may be arbitrary; the proof uses only their incidence with the recovery sets. Moreover, recovery sets need be disjoint only for requests served in the same group. We next replace the chain by a geometry with many recovery sets, then allow bounded reuse across groups.

4 The local-recovery gadget

Fix a suffix z . We encode the 2^p values $f(a, z)$, as a varies, into one evaluation-codeword. If instead we fix a code coordinate y and vary z , that coordinate becomes a shorter function $G_y(z)$. Local recovery moves across a codeword; recursion evaluates these coordinate functions down the suffix direction. [Proposition 4.2](#) records exactly what the rest of the proof needs from this construction.

Fix an integer $\ell \geq 2$ and a field $\mathbb{F}_q$ of characteristic two, where $q = 2^b$. Set

$$s_0 = \left\lfloor \frac{q-1}{\ell} \right\rfloor$$

and fix a polynomial-basis representation of $\mathbb{F}_q$. If α_i denotes the field element whose canonical b -bit representation is the integer i , choose the efficiently indexed set

$$A = \{\alpha_0, \dots, \alpha_{s_0-1}\}.$$

Let $\mathcal{V}$ be the space of polynomials

$$P(X_1, \dots, X_\ell) \in \mathbb{F}_q[X_1, \dots, X_\ell]$$

having degree less than s_0 in every variable. Repeated univariate interpolation shows that evaluation on A^ℓ uniquely determines P , and

$$\deg P \leq \ell(s_0 - 1) \leq q - 2.$$

The evaluation map

$$\text{Enc} : \mathbb{F}_q^{A^\ell} \longrightarrow \mathbb{F}_q^{\mathbb{F}_q^\ell}, \quad (P(u))_{u \in A^\ell} \longmapsto (P(y))_{y \in \mathbb{F}_q^\ell},$$

is a systematic linear code of length

$$N = q^\ell$$

and dimension

$$K = s_0^\ell = \Theta_\ell(q^\ell).$$

4.1 Packing Boolean restrictions

Split an n -bit input as

$$x = (a, z), \quad |a| = p, \quad |z| = d = n - p.$$

For each fixed suffix z , pack the 2^p bits

$$(f(a, z))_{a \in \{0,1\}^p}$$

into K field symbols, using the fixed $\mathbb{F}_2$ -basis of $\mathbb{F}_q$ and padding unused bit positions with zeros. More explicitly, interpret a prefix a as an integer $I(a) \in [2^p]_0$, and set

$$s(a) = \left\lfloor \frac{I(a)}{b} \right\rfloor, \quad j(a) = I(a) \bmod b.$$

Write $s(a)$ in base s_0 using exactly ℓ digits $i_1(a), \dots, i_\ell(a) \in [s_0]_0$, and define

$$u(a) = (\alpha_{i_1(a)}, \dots, \alpha_{i_\ell(a)}) \in A^\ell.$$

Thus bit $j(a)$ of the information symbol at $u(a)$ is $f(a, z)$. Fixed-divisor division, remainder, and base conversion have Boolean circuits of size $\text{poly}_\ell(p, \log q)$, so computing $(u(a), j(a))$ for all requests costs $\tilde{O}_\ell(t)$ gates.

Choose q , among powers of two, minimally so that

$$Kb \geq 2^p.$$

For all sufficiently large p (depending on ℓ), the corresponding code dimensions at field sizes q and $q/2$ are both $\Theta_\ell(q^\ell)$. Minimality then gives

$$q^\ell b = \Theta_\ell(2^p). \quad (2)$$

Indeed, the lower bound follows from $Kb \leq q^\ell b$. For the upper bound, the preceding candidate $q/2$ has bit width $b-1$, fails the packing inequality, and has dimension $\Omega_\ell(q^\ell)$. Hence $q^\ell(b-1) = O_\ell(2^p)$; since $b \leq 2(b-1)$ for $b \geq 2$, the upper bound in (2) follows. Equivalently,

$$\ell \log_2 q + \log_2 \log_2 q = p + O_\ell(1), \quad \log_2 q = \frac{p}{\ell} + O_\ell(\log p). \quad (3)$$

Let P_z be the polynomial whose information symbols are the packed bits. For each code coordinate $y \in \mathbb{F}_q^\ell$, define the field-valued resource

$$G_y(z) = P_z(y).$$

Each of the b coordinate bits of G_y is a Boolean function on d variables, determined by f , y , and the chosen field basis. For $r \in [b]_0$, write

$$H_{y,r}(z) = (G_y(z))_r.$$

It is useful to picture the values $P_z(y)$ as a rectangular array. A row, with z fixed and y varying, is a codeword that contains all requested prefix values. A column, with y fixed and z varying, is the resource function G_y on d variables. The line decoder reads selected entries from one row, while the recursive circuit evaluates the required columns.

4.2 Affine-line recovery

Lemma 4.1 (Line identity). *Let $P : \mathbb{F}_q^\ell \rightarrow \mathbb{F}_q$ be represented by a polynomial of total degree at most $q-2$. For every $u \in \mathbb{F}_q^\ell$ and every nonzero direction $v \in \mathbb{F}_q^\ell$,*

$$P(u) = - \sum_{\lambda \in \mathbb{F}_q^*} P(u + \lambda v).$$

In characteristic two, the minus sign may be omitted.

Proof. The restriction $Q(\lambda) = P(u + \lambda v)$ is a univariate polynomial of degree at most $q-2$. The sum of the constant monomial over $\mathbb{F}_q$ is $q \cdot 1 = 0$ in characteristic two. For $1 \leq j \leq q-2$, the sum of λ^j over $\mathbb{F}_q^*$ vanishes because j is not divisible by $q-1$, and the zero term contributes nothing. Applying these identities monomial by monomial gives $\sum_{\lambda \in \mathbb{F}_q} Q(\lambda) = 0$. Separating the term $\lambda = 0$ proves the claim. $\square$

Thus every affine line through an information point u supplies a recovery set consisting of its $q - 1$ non-target points. We identify two nonzero direction vectors when one is a nonzero scalar multiple of the other and write $[v]$ for the resulting projective direction. The number of such directions through a point is

$$D_q = \frac{q^\ell - 1}{q - 1} = 1 + q + \cdots + q^{\ell-1}, \quad q^{\ell-1} \leq D_q < 2q^{\ell-1}. \quad (4)$$

For fixed ℓ and $p \rightarrow \infty$, [Equations \(3\)](#) and [\(4\)](#) give the parameter dictionary

$$q^\ell = 2^{p+o(p)}, \quad q = 2^{p/\ell+o(p)}, \quad D_q = q^{\ell-1+o(1)} = 2^{((\ell-1)/\ell)p+o(p)}. \quad (5)$$

Two distinct lines through the same target intersect only at that target, which is absent from both recovery sets. Repeated requests for one information symbol therefore cause no additional difficulty.

The facts needed later can now be packaged without reference to the internal polynomial representation.

Proposition 4.2 (Local-recovery gadget). *Fix $\ell \geq 2$ and a split $n = p + d$, with p sufficiently large depending on ℓ , and choose $q = 2^b$ by the packing rule above. For every $f : \{0, 1\}^n \rightarrow \{0, 1\}$ there are $q^\ell b$ Boolean resource functions $H_{y,r} : \{0, 1\}^d \rightarrow \{0, 1\}$ with the following properties.*

- (i) *The resource count satisfies $q^\ell b = \Theta_\ell(2^p)$.*
- (ii) *A prefix a determines an information point $u(a) \in A^\ell$ and a bit position $j(a) \in [b]_0$ by a circuit of size $\text{poly}_\ell(p, \log q)$.*
- (iii) *For every projective direction $[v]$, the set*

$$R(u(a), [v]) = \{u(a) + \lambda v : \lambda \in \mathbb{F}_q^*\}$$

has $q - 1$ coordinates and recovers the requested bit by

$$f(a, z) = \bigoplus_{y \in R(u(a), [v])} H_{y, j(a)}(z).$$

There are $D_q = (q^\ell - 1)/(q - 1)$ such recovery sets.

- (iv) *If $U \subseteq \mathbb{F}_q^\ell$ is already occupied, at most $|U \setminus \{u(a)\}|$ directions have $R(u(a), [v]) \cap U \neq \emptyset$.*

Proof. The packing construction gives (i) and (ii). The line identity, followed by selection of coordinate bit $j(a)$ in the fixed $\mathbb{F}_2$ basis, gives (iii). For (iv), every $y \in U \setminus \{u(a)\}$ lies on exactly one projective line through $u(a)$, while the omitted target $u(a)$ blocks none. $\square$

Thus one desired bit has become $q - 1$ requests to shorter resource functions, but those coordinates may be chosen from any of D_q recovery sets. The next section turns the blocking guarantee in (iv) into a deterministic circuit and then routes only the coordinates actually selected.

5 Deterministic scheduling and routing

5.1 Disjoint scheduling inside one group

The geometric count says that a suitable line exists, but a circuit must find one from the target points. We do this greedily. Previously occupied points rule out directions; projective ranking and sorting let an ordinary Boolean circuit choose the first direction that remains.

Consider a multiset of g requested information points $u_1, \dots, u_g \in A^\ell$. Process the requests in order. Suppose disjoint recovery sets have been selected for requests $1, \dots, j-1$, and let U be their union. Then

$$|U| = (j-1)(q-1) < gq.$$

For the target u_j , every point $y \in U \setminus \{u_j\}$ forbids at most one projective direction, namely the direction of $y - u_j$. The point u_j itself forbids none because it is omitted from the new recovery set.

Lemma 5.1 (Projective indexing). *Projective directions in $\mathbb{F}_q^\ell$ have rank and unrank circuits of size $\text{poly}_\ell(b)$, with ranks in $[D_q]_0$.*

Implementation idea. The indexing normalizes the first nonzero coordinate to one and orders the normalized vectors by the position of that coordinate and then by their remaining base- q digits. The explicit circuits are given in [Section C.1](#).

Lemma 5.2 (Greedy line scheduler). *If*

$$gq < D_q, \tag{6}$$

then every multiset of g target points admits pairwise disjoint affine-line recovery sets. Given the targets, the recovery sets can be selected by a Boolean circuit of size

$$\tilde{O}_\ell(g^2q).$$

Proof idea. At stage j , the blocking property in [Proposition 4.2](#) leaves a valid direction because fewer than D_q directions are forbidden. The circuit ranks the forbidden directions, sorts those ranks, and selects the least missing one; stage j costs $\tilde{O}_\ell(jq)$. The treatment of duplicate targets, sentinel ranks, and the least-missing-rank circuit is given in [Section C.2](#). Summing over the stages gives the stated cost.

5.2 Request groups

Partition the t requests into C fixed groups, each of size at most

$$g = \left\lceil \frac{t}{C} \right\rceil.$$

Apply [Lemma 5.2](#) independently inside each group. A code coordinate is used at most once per group and therefore at most C times overall. The inequalities $\lceil t/C \rceil \leq t/C + 1$ and $2t \leq t^2/C + C$, the latter being the arithmetic–geometric mean inequality, show that the total scheduling cost is

$$C \cdot \tilde{O}_\ell(g^2q) = \tilde{O}_\ell\left(\frac{t^2q}{C} + Cq\right), \tag{7}$$

The construction is valid whenever (6) holds for $g = \lceil t/C \rceil$, that is, whenever

$$\left\lceil \frac{t}{C} \right\rceil q < D_q. \quad (8)$$

The parameter C exposes the central tradeoff. More groups mean fewer requests per group, making disjoint scheduling easier and cheaper. But they also permit each resource coordinate to receive more suffixes, making the recursive C -copy problem harder. The recursion works by finding an exponential choice of C that satisfies both demands. In the proof below, one encoded block supplies the room inside a group and the remaining blocks supply the capacity across groups.

5.3 Scatter and gather

Every non-target point of a selected line produces an incidence record containing a request identifier, a group, a line position, the point's code coordinate y , and the suffix to be evaluated. There are exactly $t(q-1) < tq$ such records. Inside one group, each coordinate appears in at most one record.

The routing consists of four fixed sorting-and-local-update passes. Here “key” means (group, y) ; any unlisted tie is broken by the record's original position.

Pass	Records	Sort order	Local result
Scatter match	incidence, slot	(key, type); incidence before slot	a matched slot takes its predecessor's suffix; otherwise it takes a dummy
Scatter order	incidence, filled slot	(type, key); slots first	slots are in canonical (group, y) order
Gather match	value, saved incidence	(key, type); value before incidence	a matched incidence takes its predecessor's value
Gather order	value, updated incidence	(type, request, line position); incidences first	decoder inputs are in canonical request order

Scatter records carry a d -bit suffix. After resource evaluation, value records carry the b -bit field value $G_y(z)$. The gather passes are payload-agnostic for any payload whose width is polynomial in the displayed parameters.

Lemma 5.3 (Record routing). *The suffixes can be scattered into slots indexed by $(\text{group}, y) \in [C]_0 \times \mathbb{F}_q^\ell$, and the resulting resource values can be gathered back to their request records, using circuits of total size*

$$\tilde{O}_\ell(tq + Cq^\ell).$$

Implementation idea. The four passes in the table use a constant number of sorting networks on $O(tq + Cq^\ell)$ records of polynomial width. The exact fixed-wire treatment of unmatched records and array boundaries is given in [Section C.3](#).

The lemma is the reason no tq^ℓ crossbar appears. Only the actual incidences and one dummy record per resource slot are routed.

6 The composition bound

Bounded-demand invariant. Fix f and the local-recovery gadget above. Each $H_{y,r}$ is one fixed Boolean function on d variables. Inside each request group, disjoint recovery ensures that coordinate y occurs at most once; hence the pair (y, r) receives at most one suffix from each of the C groups. Place those suffixes in the canonical group slots and pad unused slots with a fixed dummy suffix. One circuit for $H_{y,r}^{\times C}$, of size at most $L_C(d)$, serves exactly this bounded demand. There are $q^\ell b$ such pairs, and no sharing between different pairs is assumed.

Proposition 6.1 (Composition bound). *Fix $\ell \geq 2$. Let $n = p + d$ with p sufficiently large depending on ℓ , choose $q = 2^b$ minimally so that $s_0^\ell b \geq 2^p$, and suppose Equation (8) holds. Then every $f : \{0, 1\}^n \rightarrow \{0, 1\}$ satisfies*

$$C(f^{\times t}) \leq \underbrace{q^\ell b L_C(d)}_{\text{recursive resource evaluation}} + \tilde{O}_\ell \left(\underbrace{\frac{t^2 q}{C}}_{\text{line scheduling}} + \underbrace{tq}_{\text{incidences and decoding}} + \underbrace{Cq^\ell}_{\text{resource slots}} \right). \quad (9)$$

The first term is recursive: there are $q^\ell b$ Boolean resource functions, each evaluated on at most C suffixes of length d . The other three terms account for scheduling, the selected line incidences, and the padded resource slots. Increasing C makes the groups smaller, but gives every resource more suffixes to process. The induction chooses C so that the recursive term stays at scale $2^n/n$ and all three bookkeeping terms are exponentially smaller.

Proof. Map every requested prefix a to its information point $u(a)$ and bit position $j(a)$ using the canonical packing circuit. In each request group, use Lemma 5.2 to select and enumerate a recovery line. Together these steps cost $\tilde{O}_\ell(t) + \tilde{O}_\ell(t^2 q/C + Cq)$ gates. Each non-target point y on a selected line requests G_y on the suffix of the corresponding output instance.

For each pair (y, r) , use the $H_{y,r}^{\times C}$ circuit from the bounded-demand invariant on the canonical group slots. The definition of $L_C(d)$ applies separately to each induced function $H_{y,r}$, so these circuits cost $q^\ell b L_C(d)$ in total. The complete scatter, gather, and fixed-wire compaction pipeline costs $\tilde{O}_\ell(tq + Cq^\ell)$ by Lemma 5.3. Field addition is coordinatewise XOR in the fixed $\mathbb{F}_2$ basis, so XORing the $q - 1$ returned b -bit values of a request and selecting bit $j(a)$ costs $O(qb)$ per request and $O(tqb)$ in total. Since each selected line passes through its target $u(a)$ and $\deg P_z \leq q - 2$, Lemma 4.1 shows that this XOR equals $P_z(u(a))$ for the request's suffix z , and bit $j(a)$ of that symbol is $f(a, z)$ by the packing of Section 4.1. The circuit therefore computes $f^{\times t}$.

The polynomial P_z and the functions G_y specify the nonuniform resource truth tables; no run-time encoder is being assumed. Since $q \geq 2$, all smaller terms in this gate ledger are absorbed by the overhead in (9), proving the bound. $\square$

7 The equal-block induction

We now apply the composition bound recursively. At level k , split the input into $k+1$ approximately equal blocks, encode one block, and invoke the preceding level on the remaining k .

For an integer $k \geq 1$, let P_k denote the assertion that, for every fixed $0 \leq \gamma < k/(k+1)$, there is a constant $A_{k,\gamma}$ such that, for all sufficiently large n , every $f : \{0,1\}^n \rightarrow \{0,1\}$, and every integer $1 \leq t \leq 2^{\gamma n}$,

$$C(f^{\times t}) \leq A_{k,\gamma} \frac{2^n}{n}.$$

Every comparison below has a fixed positive linear margin in its base-two logarithm. We use without further comment that such a margin absorbs the $o(n)$ terms, every fixed polynomial factor hidden by $\tilde{O}$, and the $O(1)$ changes caused by floors and ceilings.

7.1 Two equal blocks: the base case

Lemma 7.1 (Two-block base). P_1 holds.

Here $C = 1$, so no mass-production hypothesis is used inside a resource. The two equal blocks make the resource term proportional to $2^{n/2}L(n/2)$; the one-half threshold comes from the quadratic scheduler term t^2q .

Proof. Fix $0 \leq \gamma < 1/2$, let $\eta = 1/2 - \gamma > 0$, and suppose $t \leq 2^{\gamma n}$. Put

$$m = \left\lfloor \frac{n}{2} \right\rfloor, \quad p = m, \quad d = n - m, \quad C = 1.$$

Thus $n = 2m + O(1)$ and all recovery sets will be globally disjoint. In block units, the request and code parameters are

$$\log_2 t \leq m - \eta n + O(1), \quad \log_2 q = \frac{m}{\ell} + o(n), \quad \log_2 D_q = \left(1 - \frac{1}{\ell}\right)m + o(n).$$

Choose a sufficiently large fixed ℓ so that

$$\frac{1}{\ell} < \eta. \tag{10}$$

Since $2m/\ell \leq n/\ell < \eta n$, this gives

$$m - \eta n + \frac{m}{\ell} < m - \frac{m}{\ell}.$$

The strict linear margin now implies $tq < D_q$ for all sufficiently large n .

The resource term in Equation (9) is, by Shannon–Lupanov synthesis,

$$q^\ell bL(d) = O_\ell \left(2^p \frac{2^d}{d} \right) = O_\gamma(2^n/n).$$

The scheduler exponent is

$$\log_2(t^2q) \leq 2m - 2\eta n + \frac{m}{\ell} + o(n) < n,$$

where the strict inequality follows from (10). Likewise, $\log_2(tq) \leq m - \eta n + m/\ell + o(n) < n$ and $\log_2 q^\ell = m + o(n) < n$. After the suppressed polynomial factors are restored, every overhead term is $2^{(1-\Omega_\gamma(1))n}$ and is negligible compared with $2^n/n$. The same strict margins absorb all rounding. $\square$

7.2 Adding one equal block

Lemma 7.2 (Equal-block step). *For every integer $k \geq 2$, P_{k-1} implies P_k .*

Proof. Fix $0 \leq \gamma < k/(k+1)$. If $\gamma < (k-1)/k$, the assertion is already P_{k-1} . We may therefore assume

$$\frac{k-1}{k} \leq \gamma < \frac{k}{k+1}.$$

Let

$$\eta = \frac{k}{k+1} - \gamma > 0, \quad m = \left\lfloor \frac{n}{k+1} \right\rfloor, \quad p = m, \quad d = n - m.$$

Write $n = (k+1)m + r$ with $0 \leq r \leq k$, and choose the number of request groups to be

$$C = \left\lceil 2^{(k-1)m - \eta n/2} \right\rceil. \quad (11)$$

This choice splits the available slack in half. Dividing the requests among C groups puts each group $\eta n/2$ below the one-block exponent m ; measured on the remaining d bits, the exponent of C is also a fixed amount below the threshold supplied by P_{k-1} . The assumption $\gamma \geq (k-1)/k$ ensures that the exponent in (11) is positive for all sufficiently large n .

Let $\beta = (k-1)/k$. Since $d = km + r$,

$$\beta d - \log_2 C = \frac{\eta n}{2} + \frac{k-1}{k}r + O(1).$$

In particular, $C \leq 2^{(\beta - \eta/4)d}$ for all sufficiently large n . The exponent $\beta - \eta/4$ is fixed and strictly below the threshold for P_{k-1} , so the inductive hypothesis gives

$$L_C(d) = O_{k,\gamma}(2^d/d).$$

By Equation (2), the resource term is therefore

$$q^\ell b L_C(d) = O_{k,\gamma,\ell} \left(2^m \frac{2^d}{d} \right) = O_{k,\gamma,\ell}(2^n/n).$$

The remaining estimates are easiest to read in block units. Since $\log_2 t \leq \gamma n = km - \eta n + O_k(1)$, the size $g = \lceil t/C \rceil$ of a request group satisfies

$$\log_2 g \leq m - \frac{\eta n}{2} + O_k(1).$$

The code parameters satisfy

$$\log_2 q = \frac{m}{\ell} + o(n), \quad \log_2 D_q = \left(1 - \frac{1}{\ell}\right)m + o(n).$$

Choose a sufficiently large fixed ℓ so that

$$\frac{2}{(k+1)\ell} < \frac{\eta}{2}. \quad (12)$$

Then $2m/\ell < \eta n/2$, and hence

$$m - \frac{\eta n}{2} + \frac{m}{\ell} < m - \frac{m}{\ell}.$$

Together with (4), this is exactly the strict exponent margin required for Equation (8).

Finally, the complete overhead ledger in Equation (9) is

$$\begin{aligned}\log_2\left(\frac{t^2q}{C}\right) &\leq (k+1)m - \frac{3\eta n}{2} + \frac{m}{\ell} + o(n) < n, \\ \log_2(tq) &\leq n - m - \eta n + \frac{m}{\ell} + o(n) < n, \\ \log_2(Cq^\ell) &\leq n - m - \frac{\eta n}{2} + o(n) < n.\end{aligned}$$

The first strict inequality uses $(k+1)m \leq n$ and (12); the other two use $m = \Theta_k(n)$ and $\ell > 1$. Thus every overhead term is exponentially smaller than $2^n/n$. This proves P_k . $\square$

Proof of Theorem 1.1. By Lemmas 7.1 and 7.2, P_k holds for every $k \geq 1$. Fix $0 \leq \gamma < 1$ and choose k so that $\gamma < k/(k+1)$. Then P_k gives the asserted bound beyond a threshold N_γ . For the finitely many integers $1 \leq n < N_\gamma$, use $C(f^{\times t}) \leq tC(f)$ and enlarge A_γ to dominate $ntC(f)/2^n$ over all relevant n , f , and t . This proves the statement for every $n \geq 1$. $\square$

8 Implications and open problems

The theorem settles the fixed-rate question: every fixed $\gamma < 1$ is achievable at the optimal worst-case order. It leaves three quantitative issues open: whether recursion is necessary, whether the one-copy leading constant can be retained, and where the plateau ends as the copy exponent approaches one.

Can recursion be removed? The disjoint base case isolates the reason for the recursion. With $C = 1$ and $d = \delta n$, the resource term remains $O(2^n/(\delta n))$, and for large fixed ℓ there are enough directions whenever $\gamma < 1 - \delta$. It is the quadratic scheduling term t^2q , not line availability, that restricts this argument to $\gamma < 1/2$. Grouping divides that term by C , while recursion pays for the resulting C simultaneous uses of each resource. At the exponent level, this is the update $\beta \mapsto 1/(2 - \beta)$ underlying the equal-block sequence. A disjoint-line scheduler of size $\tilde{O}(tq)$ would remove the need for recursion: for any fixed $\gamma < 1$, one could choose $0 < \delta < 1 - \gamma$ and take $C = 1$. Finding such a scheduler, or ruling one out, is therefore a natural next problem.

Can the leading constant be sharpened? The elementary proof loses a constant in two places. Its product code has rate bounded away from one, and its resource functions are synthesized on a shorter suffix, with denominator d instead of n . Scheduler and routing costs are exponentially lower order, so optimizing those circuits cannot repair either loss.

Section A removes the code-rate loss with a rate- $1 - o(1)$ lifted Reed–Solomon code, absorbs field-size rounding into a subexponential number of codewords, and then chooses unequal block lengths. The result is

$$\Lambda_\gamma \leq \frac{\Lambda_0}{1 - \gamma} = \frac{1}{1 - \gamma}.$$

The factor $1/(1 - \gamma)$ is exactly the suffix-denominator loss in this accounting. We conjecture that unrestricted circuits avoid it and satisfy $\Lambda_\gamma = \Lambda_0 = 1$. The appendix states an equivalent transversal-synthesis problem and proves a structural constraint on any such improvement.

What changes if errors are allowed? For a fixed error density $0 < \rho < 1/2$, worst-case approximate synthesis still has order $2^n/n$, but its sharp one-copy coefficient is the rate-distortion factor $1 - H_2(\rho)$, where H_2 is binary entropy [11]; the more general dependence on the approximation advantage and on partial domains was determined in [2]. Section B shows that this lossy front end is compatible with our local-recovery construction. If one first chooses a fixed function g that disagrees with f on at most a ρ -fraction of the domain and then computes g exactly on every copy, the resulting mass-production coefficient is at most

$$\frac{1 - H_2(\rho)}{1 - \gamma}.$$

A weaker model may permit a more substantial simplification. If only the expected fraction of wrong coordinates on a uniformly random input tuple is bounded, one can contemplate assigning recovery directions in advance, sorting the resulting incidences, and abandoning requests whose lines collide. Two fixed-direction affine lines through independent uniform targets intersect with probability on the order of $q^{2-\ell}$, suggesting that almost all requests could be served with demand one and hence without recursion. Such a circuit need not compute one fixed approximant g : the answer at one coordinate can depend on the other inputs in its batch. We therefore leave this batch-dependent error model as a separate problem rather than fold it into the theorem below.

Where does the plateau end? The theorem reaches $2^{(1-\varepsilon)n}$ copies for every fixed $\varepsilon > 0$, but does not treat a rate that approaches one with n . Some degradation in that regime is forced. If $n \geq 2$ and f depends on all n of its variables, every output of a circuit for $f^{\times t}$ is a gate, and each of the tn inputs must feed at least one gate, since otherwise the corresponding output would not depend on it. Since gates have fan-in at most two,

$$C(f^{\times t}) \geq \frac{tn}{2}.$$

A bound of the form $A 2^n/n$ therefore cannot hold once $t > 2A \cdot 2^n/n^2$; in particular, a uniform $O(2^n/n)$ plateau must fail for $t = \omega(2^n/n^2)$. The theorem still leaves the detailed mass-production curve throughout the near-endpoint regime $t = 2^{n-o(n)}$ unresolved. Determining how far the plateau extends, and then characterizing the growth beyond it, would require a quantitatively sharper schedule or a different recursive decomposition.

Beyond unrestricted circuits. The construction uses arbitrary routing, finite-field linear reconstruction, and free fan-out. Formula size, bounded-fan-out circuits, monotone circuits, bounded-depth circuits, and branching programs may have different mass-production curves. In those settings, simultaneous demand may become a source of genuine direct-sum lower bounds. Even in the unrestricted model, affine lines are not the only possible local-recovery geometry. Batch codes, locally recoverable codes, or higher-dimensional varieties may improve the group-size versus decoding-cost tradeoff and reduce the number of recursive levels.

A Leading-coefficient bounds

The constant A_γ in [Theorem 1.1](#) also covers finitely many small input lengths, so it is not a meaningful leading coefficient. This appendix tracks the asymptotic coefficient instead. The first subsection records the exact, but large, constant from the elementary product code. We then replace that code by a high-rate lifted code and choose unequal block lengths, which reduces the coefficient to $1/(1 - \gamma)$. Readers interested only in that final bound may skip to [Proposition A.3](#).

Define the one-copy and mass-production coefficients

$$\Lambda_0 = \limsup_{n \rightarrow \infty} \frac{nL(n)}{2^n}$$

and, for $0 < \gamma < 1$,

$$\Lambda_\gamma = \limsup_{n \rightarrow \infty} \max_{\substack{f: \{0,1\}^n \rightarrow \{0,1\} \\ 1 \leq t \leq 2^{\gamma n}}} \frac{nC(f^{\times t})}{2^n}.$$

Thus [Theorem 1.1](#) says in particular that $\Lambda_\gamma < \infty$. The matched Shannon–Lupanov bounds of Frandsen and Miltersen [5] give, in our basis,

$$\Lambda_0 = 1. \tag{13}$$

A.1 The elementary product-code ledger

We first isolate the constant lost in packing. For fixed $\ell \geq 2$ and prefix length p , let $q = 2^b$ be minimal subject to

$$\left\lfloor \frac{q-1}{\ell} \right\rfloor^\ell b \geq 2^p,$$

as in [Section 4.1](#), put $s_0 = \lfloor (q-1)/\ell \rfloor$, and write

$$R_\ell(p) = \frac{q^\ell b}{2^p}, \quad \rho_\ell = \limsup_{p \rightarrow \infty} R_\ell(p).$$

The code rate and the restriction to fields of power-of-two size give

$$\ell^\ell \leq \rho_\ell \leq (2\ell)^\ell. \tag{14}$$

Indeed, the packing inequality implies $R_\ell(p) \geq (q/s_0)^\ell \rightarrow \ell^\ell$. For the reverse inequality, the preceding field size $q/2$ fails the packing condition. If $s'_0 = \lfloor (q/2 - 1)/\ell \rfloor$, then

$$2^p > (s'_0)^\ell (b-1), \quad R_\ell(p) < \left(\frac{q}{s'_0} \right)^\ell \frac{b}{b-1} \rightarrow (2\ell)^\ell.$$

We can now state exactly what the elementary construction proves.

Proposition A.1 (Leading coefficient from the equal-block construction). *For $0 < \gamma < 1$, put*

$$S_k = k - (k+1)\gamma.$$

Then

$$\Lambda_\gamma \leq \Lambda_0 \inf_{\substack{k \geq 1 \\ S_k > 0}} (k+1) \inf_{\substack{\ell_1, \dots, \ell_k \in \mathbb{Z}_{\geq 3} \\ \sum_{j=1}^k 2/\ell_j < S_k}} \prod_{j=1}^k \rho_{\ell_j}. \tag{15}$$

Both infima range over nonempty discrete sets.

Proof. Fix admissible k and $\ell_1, \dots, \ell_k$. Since

$$\sum_{j=1}^k \frac{2}{\ell_j} < S_k < k,$$

we may choose numbers Δ_j such that

$$\frac{2}{\ell_j} < \Delta_j < 1, \quad \sum_{j=1}^k \Delta_j = S_k. \quad (16)$$

The Δ_j allocate the total exponent slack S_k among the k levels. The lower bound $2/\ell_j$ is the amount needed to fit a recovery line inside the available directions at level j . Set $\eta_0 = 0$ and, for $1 \leq j \leq k$, set

$$\eta_j = \frac{1}{j+1} \sum_{i=1}^j \Delta_i, \quad \gamma_j = \frac{j}{j+1} - \eta_j.$$

Then $\gamma_0 = 0$, $\gamma_k = \gamma$, and

$$(j+1)\eta_j - j\eta_{j-1} = \Delta_j. \quad (17)$$

Because every $\Delta_i < 1$, we also have $\gamma_j > 0$ for $j \geq 1$.

Build the construction from level 0 outward. At level j , the current input length is $n_j = (j+1)m + O_j(1)$, and the target copy count is $T_j = \lceil 2^{\gamma_j n_j} \rceil$. Use one block of length m for the code, leaving a suffix of length $d_j = jm + O_j(1)$. Choose

$$C_j = \lfloor 2^{\gamma_{j-1} d_j} \rfloor$$

request groups. The recursive resource call is therefore within the copy range already constructed at level $j-1$. The logarithm of C_j is

$$\gamma_{j-1} jm + o(m) = (j-1 - j\eta_{j-1})m + o(m).$$

The logarithm of the resulting group size $\lceil T_j/C_j \rceil$ is, by (17),

$$(\gamma_j(j+1) - \gamma_{j-1}j)m + o(m) = (1 - \Delta_j)m + o(m).$$

Thus $\Delta_j m$ is literally the slack between a request group and the one-block scale 2^m . Since $\log_2 q = m/\ell_j + o(m)$ and $\log_2 D_q = (1 - 1/\ell_j)m + o(m)$, the inequality $\Delta_j > 2/\ell_j$ is precisely the strict exponent margin needed for the disjoint-line schedule.

For completeness, put $c_j = j - 1 - j\eta_{j-1}$, the exponent of C_j in block units. Ignoring lower-order terms, the three bookkeeping terms in Equation (9) have exponents

$$\begin{aligned} \frac{T_j^2 q}{C_j} : \quad & c_j + 2(1 - \Delta_j) + \frac{1}{\ell_j} < j + 1, \\ T_j q : \quad & c_j + (1 - \Delta_j) + \frac{1}{\ell_j} < j + 1, \\ C_j q^{\ell_j} : \quad & c_j + 1 < j + 1. \end{aligned}$$

Here $c_j \leq j - 1$, and the first two strict inequalities follow from $\Delta_j > 2/\ell_j$. Each term is therefore exponentially smaller than the local scale $2^{(j+1)m}/m$. The number of levels is fixed, so these terms remain lower order after the outer resource factors are restored.

Only the recursive resource term contributes to the leading coefficient. The packing loss at level j is at most ρ_{ℓ_j} in the limit, and changing the denominator from the remaining jm variables to the full $(j+1)m$ variables contributes $(j+1)/j$. Consequently the coefficients obtained after successive levels satisfy

$$B_j \leq \frac{j+1}{j} \rho_{\ell_j} B_{j-1}, \quad B_0 = \Lambda_0.$$

Iteration telescopes the normalization factors and gives

$$B_k \leq \Lambda_0(k+1) \prod_{j=1}^k \rho_{\ell_j}.$$

Taking the two infima proves (15). Floors, ceilings, and fixed polynomial factors do not affect a limsup leading coefficient because every inequality in (16) is strict. $\square$

There is a simple near-closed-form choice. For every k with $S_k > 0$, put

$$\delta_k = \frac{S_k}{k} = 1 - \gamma - \frac{\gamma}{k}, \quad \ell_k = \left\lfloor \frac{2}{\delta_k} \right\rfloor + 1. \quad (18)$$

Taking every $\Delta_j = \delta_k$ and every code parameter equal to ℓ_k in Proposition A.1, then using (14), gives

$$\Lambda_\gamma \leq \Lambda_0(k+1) \rho_{\ell_k}^k \leq \Lambda_0(k+1) (2\ell_k)^{k\ell_k}. \quad (19)$$

Equal slack also minimizes the continuous relaxation. Under the estimates $\rho_\ell \leq (2\ell)^\ell$ and $\ell = 2/\Delta$, the logarithmic cost of a level is

$$\psi(\Delta) = \frac{2}{\Delta} \log_2 \frac{4}{\Delta}.$$

The function ψ is strictly convex on $(0, 1)$, so Jensen's inequality shows that equal Δ_j minimize the relaxed cost for fixed k . The choice in (18) implements this continuous optimizer up to integer rounding.

Finally, write $\varepsilon = 1 - \gamma$ and let

$$k = \left\lceil \frac{2\gamma}{\varepsilon} \right\rceil.$$

As $\varepsilon \downarrow 0$, this choice has

$$k = \frac{2 + o(1)}{\varepsilon}, \quad \delta_k = \left(\frac{1}{2} + o(1) \right) \varepsilon, \quad \ell_k = \frac{4 + o(1)}{\varepsilon}.$$

Equation (19) therefore yields

$$\Lambda_\gamma \leq \Lambda_0 2^{(8+o(1))\varepsilon^{-2} \log_2(1/\varepsilon)}. \quad (20)$$

Within the same continuous relaxation, writing $k = a/\varepsilon$ makes $k\ell_k \sim 2a^2/((a-1)\varepsilon^2)$; this is minimized at $a = 2$. Thus the choices above optimize the leading logarithmic term of the equal-block bound. The remaining losses in (20) come from the evaluation-code rate and power-of-two field rounding. A sharper scheduler alone does not remove those losses, and (15) is not a lower bound on the true value of Λ_γ .

A.2 Removing the packing loss

Two separate devices remove the packing loss. Lifted Reed–Solomon codes retain the degree bound on every affine line while their rate tends to one. We also use a subexponential number of codewords, which lets us choose the field just below the desired size instead of rounding one codeword up to the next power of two. Neither change alters the line scheduler.

Lemma A.2 (Asymptotically lossless local recovery). *Fix $\ell \geq 3$. For every sufficiently large p , there exist a power of two $q = 2^b$, an integer B , and a family of $Bq^\ell b$ Boolean resource functions with the recovery and blocking properties of [Proposition 4.2](#), except that a resource coordinate also carries a codeword index in $[B]_0$. As $p \rightarrow \infty$, these parameters satisfy*

$$\log_2 B = o(p), \quad Bq^\ell b = (1 + o_p(1))2^p, \quad \log_2 q = \frac{p}{\ell} - o(p). \quad (21)$$

If $K_{\ell,q}$ denotes the dimension of the code over $\mathbb{F}_q$, the resulting composition bound is

$$\begin{aligned} C(f^{\times t}) &\leq (1 + o_p(1))2^p L_C(d) \\ &\quad + \tilde{O}_\ell \left(K_{\ell,q} + t + \frac{t^2 q}{C} + tq + CBq^\ell \right), \end{aligned} \quad (22)$$

whenever $\lceil t/C \rceil q < D_q$.

Proof. Put

$$b = \left\lfloor \frac{p - \sqrt{p}}{\ell} \right\rfloor, \quad q = 2^b, \quad N = q^\ell,$$

and use the $[\ell, q - 1, q]$ lifted Reed–Solomon code. Its length is N , every restriction to an affine line is represented by a polynomial of degree at most $q - 2$, and, for fixed ℓ , its dimension satisfies

$$K_{\ell,q} = (1 - o_q(1))N \quad (23)$$

by [\[6, Theorems 1 and 2\]](#). The code is linear, so a fixed information set makes its encoder systematic.

Let

$$B = \left\lceil \frac{2^p}{K_{\ell,q} b} \right\rceil$$

and pack the 2^p truth-table bits consecutively into the information symbols of these B codewords. The role of B is only to absorb the unused capacity in the last codeword: the chosen field makes one codeword too small, but by a subexponential rather than a constant fraction. Now

$$2^p \leq BNb \leq \frac{N}{K_{\ell,q}} 2^p + Nb = (1 + o_p(1))2^p.$$

Here [\(23\)](#) handles the first term, while $Nb/2^p \leq p2^{-\sqrt{p}}$ handles the rounding term. The same estimates give $B = 2^{o(p)}$ and the second part of [\(21\)](#).

A requested prefix determines its codeword index, its information-symbol rank, and its bit position by division by the fixed integers $K_{\ell,q}b$ and b . The information-symbol rank must then be translated to the corresponding point of the fixed information set. This is a static lookup: sort the $K_{\ell,q}$ hard-wired table records together with the query records, propagate the table value along each equal-key run, and restore query order. The circuit conventions give cost $\tilde{O}_\ell(K_{\ell,q} + t)$.

On every line, [Lemma 4.1](#) applies verbatim. Within a request group, run the same fixed greedy circuit as in [Lemma 5.2](#), now treating an occupied point as blocking only when its codeword index matches that of the current target. At stage j there are still at most $(j-1)(q-1)$ relevant point records to inspect, so the circuit size and scheduling condition are unchanged. Equivalently, if the codeword-part sizes are $g_1, \dots, g_B$, the work charged by the greedy comparisons is bounded by $\sum_\nu g_\nu^2 q \leq (\sum_\nu g_\nu)^2 q$. Coordinates in different codewords are distinct resources. Adding the codeword index to every routing key changes the slot universe from Cq^ℓ to CBq^ℓ and changes no other incidence count. There are BNb Boolean resource functions, giving the first term of [\(22\)](#); the scheduler, routing, and line decoder give the remaining terms exactly as in [Proposition 6.1](#). $\square$

A.3 Optimizing the block lengths

Rate-one packing leaves only the denominator lost when a resource on a shorter suffix is synthesized separately. Unequal blocks make that loss explicit and avoid the much larger coefficient in [\(20\)](#).

Proposition A.3 (High-rate leading coefficient). *For every fixed $0 \leq \gamma < 1$,*

$$\Lambda_\gamma \leq \frac{\Lambda_0}{1-\gamma} = \frac{1}{1-\gamma}. \quad (24)$$

Proof. The case $\gamma = 0$ follows from one-copy synthesis, so suppose $0 < \gamma < 1$. Write $\varepsilon = 1 - \gamma$, and fix $0 < \alpha < \varepsilon$. Choose a fixed $\ell \geq 3$ large enough that

$$\gamma < (\gamma + \alpha) \left(1 - \frac{2}{\ell}\right), \quad (25)$$

and then choose a fixed integer k large enough that

$$\frac{\gamma + (\gamma + \alpha)/\ell}{k} < \varepsilon - \alpha. \quad (26)$$

Both choices are possible: as $\ell \rightarrow \infty$, the right side of [\(25\)](#) tends to $\gamma + \alpha > \gamma$; after ℓ is fixed, the left side of [\(26\)](#) tends to zero as $k \rightarrow \infty$. Set

$$w_0 = \varepsilon - \alpha, \quad w_j = \frac{\gamma + \alpha}{k} \quad (1 \leq j \leq k), \quad W_j = \sum_{i=0}^j w_i, \quad \tau_j = \frac{j\gamma}{k}.$$

Thus $W_k = 1$, $\tau_0 = 0$, and $\tau_k = \gamma$. Partition an n -bit input into blocks of lengths $w_0 n, w_1 n, \dots, w_k n$, with bounded rounding errors. At stage j , the current input length is $n_j = W_j n + O_k(1)$; use the block of length $p_j = w_j n + O_k(1)$ for local recovery and the preceding n_{j-1} bits as the suffix. Set the outer copy bound to $T_j = \lceil 2^{\tau_j n} \rceil$ and use $C = T_{j-1}$ request groups.

Here w_0 is the suffix left for the final one-copy synthesis, the w_j are the recovery blocks added in successive stages, and W_j is the fraction of the original input present after stage j . The number τ_j is the copy exponent reached by that stage. Thus stage j must increase the copy exponent by γ/k while spending one recovery block of size $(\gamma + \alpha)n/k$.

The logarithmic group size is

$$\log_2(T_j/T_{j-1}) = \frac{\gamma n}{k} + O(1).$$

Since $\log_2 q = p_j/\ell - o(n)$ and $\log_2(D_q/q) = (1 - 2/\ell)p_j - o(n)$, (25) is precisely the inequality

$$\frac{\gamma}{k} + \frac{w_j}{\ell} < \left(1 - \frac{1}{\ell}\right) w_j$$

up to lower-order terms. Its left side is the exponent of a group multiplied by q ; its right side is the exponent of D_q . This verifies the scheduling condition.

It remains to check that building and routing the schedule is lower order. The slack between the current input exponent and copy exponent is

$$W_j - \tau_j = \varepsilon - \alpha + \frac{j\alpha}{k} \geq \varepsilon - \alpha. \quad (27)$$

Dividing base-two logarithms by n , the four nonrecursive contributions in (22) have exponents at most

$$\begin{aligned} \text{scheduler: } & 2\tau_j - \tau_{j-1} + \frac{w_j}{\ell} + o(1) = \tau_j + \frac{\gamma + (\gamma + \alpha)/\ell}{k} + o(1), \\ \text{incidences: } & \tau_j + \frac{w_j}{\ell} + o(1), \\ \text{slots: } & \tau_{j-1} + w_j + o(1), \\ \text{lookup: } & \max\{w_j, \tau_j\} + o(1). \end{aligned}$$

The scheduler exponent is below W_j by (26) and (27). The same two inequalities give $\tau_j + w_j/\ell < W_j$ for the incidence term. The slot inequality is the previous-stage slack $W_{j-1} - \tau_{j-1} > 0$, and the lookup inequality follows from $w_j < W_j$ and (27). Hence every nonrecursive term is exponentially smaller than $2^{n_j}/n_j$.

Only the recursive coefficient remains. At stage zero, one-copy synthesis on n_0 bits has coefficient $B_0 = \Lambda_0$. If the construction through stage $j - 1$ has coefficient B_{j-1} , the resource term in (22) is

$$(1 + o(1))2^{p_j} B_{j-1} \frac{2^{n_{j-1}}}{n_{j-1}} = \left(B_{j-1} \frac{W_j}{W_{j-1}} + o(1) \right) \frac{2^{n_j}}{n_j}.$$

Consequently

$$B_k \leq \Lambda_0 \prod_{j=1}^k \frac{W_j}{W_{j-1}} = \frac{\Lambda_0}{w_0} = \frac{\Lambda_0}{\varepsilon - \alpha}.$$

The construction for T_k copies also handles every smaller t by fixing unused inputs and discarding unused outputs. Since α was arbitrary, letting $\alpha \downarrow 0$ proves (24). $\square$

A.4 The remaining one-copy gap

The factor in (24) is the best that the present *separate-resource synthesis* ledger can certify. Indeed, in any multilevel version let w_0 be the final suffix fraction, let w_j be the fraction assigned to the j th recovery block, and let Δ_j be the increase in the ambient copy exponent at that level. Put $W_j = \sum_{i=0}^j w_i$. Line capacity requires $\Delta_j < w_j(1 - 2/\ell_j) < w_j$. Since $\sum_j \Delta_j = \gamma$, summing gives $w_0 < 1 - \gamma$. Separate synthesis then telescopes the denominator changes to

$$\prod_j \frac{W_j}{W_{j-1}} = \frac{1}{w_0} > \frac{1}{1 - \gamma}.$$

This is a limitation of that accounting, not a lower bound for unrestricted circuits.

We do not know whether unrestricted circuits must pay this factor. Our conjecture is

$$\boxed{\Lambda_\gamma = \Lambda_0 = 1 \quad \text{for every fixed } 0 \leq \gamma < 1.} \quad (28)$$

The lower bound is immediate because the maximum defining Λ_γ includes $t = 1$. The missing upper bound can be isolated as a synthesis problem. Given M arbitrary functions $h_1, \dots, h_M : \{0, 1\}^d \rightarrow \{0, 1\}$, define

$$\mathsf{T}(M, d) = \max_{h_1, \dots, h_M} \mathsf{C}((z_1, \dots, z_M) \mapsto (h_1(z_1), \dots, h_M(z_M))),$$

where the z_i are disjoint d -bit inputs. The next proposition makes the reduction exact: at every fixed positive ratio, sharp transversal synthesis is equivalent to retaining the one-copy coefficient throughout the exponential range. The name “transversal” reflects that the functions vary across the outputs and each is evaluated on its own input block.

Proposition A.4 (Transversal characterization of the sharp coefficient). *The following statements are equivalent.*

- (i) $\Lambda_\gamma = \Lambda_0$ for every fixed $0 < \gamma < 1$.
- (ii) For every sequence $p, d \rightarrow \infty$ such that $p/(p+d) \rightarrow \theta \in (0, 1)$,

$$\mathsf{T}(2^p, d) \leq (\Lambda_0 + o(1)) \frac{2^{p+d}}{p+d}. \quad (29)$$

Proof. (i) *implies* (ii). Assume (i). Given $h_1, \dots, h_{2^p} : \{0, 1\}^d \rightarrow \{0, 1\}$, define one function on $p+d$ variables by

$$f(a, z) = h_a(z), \quad a \in \{0, 1\}^p,$$

after identifying $\{0, 1\}^p$ with $[2^p]_0$. Take a circuit for $f^{\times 2^p}$ and, in its copy indexed by a , fix the first p input bits to a . The resulting circuit is the required transversal circuit, so

$$\mathsf{T}(2^p, d) \leq L_{2^p}(p+d). \quad (30)$$

Choose any fixed γ with $\theta < \gamma < 1$. For all sufficiently large $p+d$, one has $2^p \leq 2^{\gamma(p+d)}$; hence (i), the definition of Λ_γ , and (30) give (29).

(ii) *implies* (i). There are two reductions. First we remove the power-of-two restriction. Suppose $\log_2 M$ and d tend to infinity in a fixed positive ratio. Write the binary expansion of M as a sum of distinct powers of two and partition an arbitrary family $h_1, \dots, h_M$ into batches of those sizes. Put $p = \lfloor \log_2 M \rfloor$. For a fixed J , call a batch large if its size is at least 2^{p-J+1} . Every large batch has logarithm $p - O(J)$, so (29) applied to these batches gives, in total,

$$(\Lambda_0 + o(1)) \frac{2^d}{\log_2 M + d} (\text{number of functions in those batches}).$$

All remaining powers in the binary expansion sum to less than $2^{p-J+1} \leq 2^{1-J} M$. Synthesizing the corresponding functions separately therefore costs at most $2^{1-J} M L(d)$. First let the parameters tend to infinity. Relative to $M 2^d / (\log_2 M + d)$, this remainder is at most

$$2^{1-J} (\Lambda_0 + o(1)) \frac{\log_2 M + d}{d},$$

Because $\log_2 M/d$ stays bounded, this quantity tends to zero when $J \rightarrow \infty$. Using (13), we obtain the unrestricted form

$$\mathsf{T}(M, d) \leq (\Lambda_0 + o(1)) \frac{M2^d}{\log_2 M + d}. \quad (31)$$

The second reduction replaces the separately synthesized leaves of the high-rate construction by one transversal circuit. Fully unroll the fixed number of levels used in the proof of Proposition A.3. Let d be the length of the suffix left at stage zero, and let M be the number of leaf resource functions, counted with multiplicity. At each level, rate-one packing multiplies the number of leaves by $(1 + o(1))2^{p_j}$, where p_j is the new recovery-block length. Since the number of levels is fixed and $\sum_j p_j = n - d + o(n)$, multiplication over the levels gives

$$M2^d = (1 + o(1))2^n, \quad \log_2 M + d = n + o(n),$$

with both $\log_2 M$ and d fixed positive fractions of n .

Before any routed suffixes are identified, give each of the M leaf occurrences its own d -bit input slot. The leaves are then an arbitrary family of M functions, each requested once, so we may apply (31) to the leaf family, and then identify input slots that carry the same routed suffix. Identification does not add gates. The leaf circuit therefore has

$$(\Lambda_0 + o(1)) \frac{2^n}{n}$$

gates. All scheduler, routing, lookup, and decoding terms in the unrolled construction are exponentially smaller, exactly as in the proof of Proposition A.3. This proves $\Lambda_\gamma \leq \Lambda_0$ for every fixed $\gamma < 1$; the reverse inequality follows from $t = 1$. Hence (i) holds. $\square$

There is a simple obstruction to proving the transversal bound by placing separate syntheses side by side. Let G be any circuit for a family $h_1(z_1), \dots, h_M(z_M)$ on disjoint input blocks. For a gate v of G , let $b(v)$ be the number of blocks on which the Boolean function computed at v depends essentially. Call G *block-respecting* if $b(v) \leq 1$ for every gate.

Proposition A.5 (Block-support incidence bound). *Every circuit G for such a family satisfies*

$$\sum_{i=1}^M \mathsf{C}(h_i) \leq \sum_{v \in \text{Gate}(G)} b(v). \quad (32)$$

Consequently, if every gate depends essentially on variables from at most w input blocks, then

$$|G| \geq \frac{1}{w} \sum_{i=1}^M \mathsf{C}(h_i).$$

In particular, the worst-case size among block-respecting circuits is exactly $ML(d)$.

Proof. Fix i and assign arbitrary constants to every input block other than z_i . The restricted circuit still computes h_i at its i th output. If i is not in the block support of a gate v , the function computed at v becomes a constant under this restriction, so that gate can be replaced by a free constant. Retaining all other gates gives a circuit for h_i with at most the following number of gates:

$$|\{v : i \text{ lies in the block support of } v\}|$$

Sum this inequality over i and interchange the two counts to obtain (32). If all block supports have size at most w , the second assertion follows. For $w = 1$, taking every h_i to have complexity $L(d)$ gives the lower bound $ML(d)$, while placing M optimal one-copy circuits side by side gives the matching upper bound. $\square$

For $M = 2^p$ and p/d tending to a fixed positive constant, the last case has normalized coefficient

$$\frac{(p+d)ML(d)}{2^{p+d}} = (\Lambda_0 + o(1)) \frac{p+d}{d}.$$

This is exactly the factor lost by separate-resource synthesis. More quantitatively, a circuit of the conjectured sharp size for a family with $C(h_i) = (\Lambda_0 - o(1))2^d/d$ for every i would have to satisfy

$$\frac{1}{|G|} \sum_{v \in \text{Gate}(G)} b(v) \geq (1 - o(1)) \frac{p+d}{d}. \quad (33)$$

Separate Lupanov circuits therefore cannot recover the missing factor. A sharp construction would need gates whose computed functions depend on at least $(p+d)/d$ input blocks on average; that foreign-block dependence must then disappear before the block-local outputs.

Ordinary multi-output Lupanov synthesis does not prove (29): its outputs share one input, whereas here both the functions and their input blocks vary. Likewise, Paul's nonadditivity examples [10] show that unrestricted circuits can share across disjoint blocks, but do not provide a worst-case synthesis theorem for an arbitrary family. Thus improving $1/(1-\gamma)$ to 1 is exactly the transversal-synthesis problem above, rather than a further optimization of the separate-resource ledger. In particular, a counterexample to (29) along any fixed ratio would refute (28).

B Approximate mass production

This appendix considers a fixed approximant rather than a batch-dependent notion of error. For two functions $f, g : \{0, 1\}^n \rightarrow \{0, 1\}$, put

$$\text{dist}(f, g) = 2^{-n} |\{x \in \{0, 1\}^n : f(x) \neq g(x)\}|,$$

and, for $0 \leq \rho < 1/2$, define

$$L_t^{(\rho)}(n) = \max_{f: \{0, 1\}^n \rightarrow \{0, 1\}} \min_{\substack{g: \{0, 1\}^n \rightarrow \{0, 1\} \\ \text{dist}(f, g) \leq \rho}} C(g^{\times t}).$$

Thus the bad set is chosen after f but is fixed before the batch is evaluated; the circuit computes the same approximating function g exactly on all t input blocks. Define the corresponding leading coefficient by

$$\Lambda_\gamma^{(\rho)} = \limsup_{n \rightarrow \infty} \max_{1 \leq t \leq 2^{\gamma n}} \frac{n L_t^{(\rho)}(n)}{2^n}.$$

Let

$$H_2(\rho) = -\rho \log_2 \rho - (1 - \rho) \log_2 (1 - \rho), \quad H_2(0) = 0,$$

and abbreviate $A_\rho = 1 - H_2(\rho)$. Pippenger's asymptotic approximation theorem gives, in the present convention,

$$\Lambda_0^{(\rho)} = A_\rho; \tag{34}$$

his lower bound holds even if every unary and binary Boolean gate is available, while the matching upper bound in our basis also follows from the construction below. See [11]. Andreev, Clementi, and Rolim later determined the optimal order uniformly over the approximation advantage and the size of a partial domain [2]. For fixed ρ , the next lemma gives the part of the one-copy rate–distortion construction needed here and makes it locally recoverable.

Lemma B.1 (Lossy block packing). *Fix $0 \leq \rho < 1/2$. For every sufficiently large power of two s , there is a set $\mathcal{K}_s \subseteq \{0, 1\}^s$ and an indexing of its elements by r_s bits such that*

$$r_s = (A_\rho + o_s(1))s \tag{35}$$

and every $v \in \{0, 1\}^s$ differs from some member of $\mathcal{K}_s$ in at most ρs coordinates. The indexed coordinate map

$$\text{Dec}_s : \{0, 1\}^{r_s} \times [s]_0 \longrightarrow \{0, 1\}$$

has a circuit of size $2^{O(s)}$.

Consequently, if s divides 2^p , then each row $a \mapsto f(a, z)$ can be replaced by a row $a \mapsto g(a, z)$ with at most $\rho 2^p$ errors and represented by

$$m_p = \frac{2^p}{s} r_s = (A_\rho + o_s(1))2^p \tag{36}$$

bits.

Proof. Let

$$V_s = \sum_{j=0}^{\lfloor \rho s \rfloor} \binom{s}{j} = 2^{H_2(\rho)s + o(s)}$$

be the size of a Hamming ball of radius $\lfloor \rho s \rfloor$. Choose

$$M_s = \left\lceil (s+1) \ln 2 \frac{2^s}{V_s} \right\rceil$$

independent uniform centers. A fixed word is uncovered with probability at most $\exp(-M_s V_s / 2^s)$, so a union bound over all 2^s words is smaller than one. Some choice of at most M_s centers therefore covers the cube. Conversely, the union of the balls about any covering set has size at most $|\mathcal{K}_s| V_s$, so every covering set has at least $2^s / V_s$ members. It follows that, after deleting duplicate centers and putting $r_s = \lceil \log_2 |\mathcal{K}_s| \rceil$, one has (35).

Fix an ordering of $\mathcal{K}_s$. On a valid index and a coordinate, Dec_s returns that coordinate of the indexed center; define it arbitrarily on unused indices. A truth-table circuit has size $2^{O(s)}$ because $r_s \leq s+1$. Finally, partition a length- 2^p row into $2^p/s$ blocks and choose a covering center separately for every block. Concatenating their indices gives (36), and the sum of their Hamming errors is at most $\rho 2^p$. $\square$

The important point is that the indices, rather than the approximated row itself, are fed into the local-recovery code. All approximation occurs before recovery; the resource functions and line decoders remain exact.

Lemma B.2 (Lossy local-recovery composition). *Fix $0 \leq \rho < 1/2$ and $\ell \geq 3$. Let $n = p + d$, and let s be a power of two with $s \rightarrow \infty$, $s = o(p)$, and $s = O(\log n)$. Suppose s divides 2^p . For every $f : \{0, 1\}^n \rightarrow \{0, 1\}$, there is a $g : \{0, 1\}^n \rightarrow \{0, 1\}$ with $\text{dist}(f, g) \leq \rho$ and parameters $q = 2^b$, B , and $K_{\ell, q}$ satisfying*

$$\log_2 q = \frac{p}{\ell} - o(p), \quad \log_2 B = o(p), \quad Bq^\ell b = (A_\rho + o(1))2^p,$$

for which

$$\begin{aligned} C(g^{\times t}) &\leq (A_\rho + o(1))2^p L_C(d) \\ &\quad + \tilde{O}_\ell \left(K_{\ell, q} + tr_s + \frac{(tr_s)^2 q}{C} + tr_s q + CBq^\ell \right) + t2^{O(s)}, \end{aligned} \tag{37}$$

whenever

$$\left\lceil \frac{tr_s}{C} \right\rceil q < D_q. \tag{38}$$

Here the $o(1)$ terms tend to zero as $p, s \rightarrow \infty$.

Proof. Apply Lemma B.1 independently to every suffix z . For each block of the row $a \mapsto f(a, z)$, choose one covering center and concatenate its r_s -bit index. This produces an m_p -bit compressed row w_z . The decoded rows define one fixed function g , and the row-wise error guarantee gives $\text{dist}(f, g) \leq \rho$.

We now repeat the construction of Lemma A.2 with the message length m_p in place of 2^p . Explicitly, put

$$b = \left\lfloor \frac{\log_2 m_p - \sqrt{p}}{\ell} \right\rfloor, \quad q = 2^b, \quad N = q^\ell,$$

use the $[\ell, q-1, q]$ lifted Reed–Solomon code of dimension $K_{\ell,q} = (1 - o_q(1))N$, and take

$$B = \left\lceil \frac{m_p}{K_{\ell,q}b} \right\rceil.$$

Because $\log_2 m_p = p + O(1)$ and $Nb/m_p \leq p2^{-\sqrt{p}+O(1)}$, the same calculation as in (21) gives

$$BNb = (1 + o(1))m_p = (A_\rho + o(1))2^p$$

and the asserted estimates for q and B .

Given a prefix a , split its bits into its block number and its within-block coordinate. For each of the r_s bits of that block's codeword index, compute the corresponding position in w_z and treat it as one local-recovery target. Thus t requested values of g create tr_s recovery targets. Apply the scheduler and the four-pass routing construction to those targets. The bounded-demand invariant is unchanged, and there are BNb Boolean resource functions, which gives the first term of (37). Replacing t by tr_s in the remaining terms of (22) gives the displayed scheduling, incidence, slot, and lookup costs. Once the r_s index bits for an output have been recovered exactly, evaluate Dec_s at the within-block coordinate. The t decoder copies cost $t2^{O(s)}$. Their outputs are exactly $g^{\times t}$. $\square$

Proposition B.3 (Approximate leading coefficient). *For every fixed $0 \leq \rho < 1/2$ and $0 \leq \gamma < 1$,*

$$\Lambda_\gamma^{(\rho)} \leq \frac{1 - H_2(\rho)}{1 - \gamma}. \quad (39)$$

Proof. First suppose $\gamma = 0$. Take $p = \lceil \sqrt{n} \rceil$, $d = n - p$, any fixed $\ell \geq 3$, $t = C = 1$, and a power of two $s = \Theta(\log n)$ with $s \mid 2^p$. Since $r_s = O(\log n)$, the scheduling condition (38) holds for large n . Every term on the second line of (37), as well as the final decoder term, is $2^{o(n)}$. Using $L(d) = (1 + o(1))2^d/d$ in the first term gives

$$L_1^{(\rho)}(n) \leq (A_\rho + o(1))2^p \frac{2^d}{d} = (A_\rho + o(1)) \frac{2^n}{n}.$$

Together with (34), this also recovers the sharp one-copy coefficient.

Now let $0 < \gamma < 1$. Use the block fractions, current lengths, and copy bounds from the proof of Proposition A.3: for arbitrary $0 < \alpha < 1 - \gamma$, they satisfy

$$w_0 = 1 - \gamma - \alpha, \quad W_k = 1, \quad n_j = W_j n + O(1), \quad T_j = \lceil 2^{j\gamma n/k} \rceil.$$

Run the exact high-rate construction through stage $k-1$. Its coefficient is at most

$$\Lambda_0 \frac{W_{k-1}}{w_0}.$$

At the final stage, apply Lemma B.2 to the outermost recovery block, with $t = T_k$ and $C = T_{k-1}$. The factor $r_s = O(\log n)$ adds only $o(n)$ to every base-two logarithm, so all strict scheduling and overhead inequalities from the exact proof remain valid. The new decoder term has exponent $\gamma + o(1) < 1$ and is also lower order.

Only the final resource coefficient changes. It is at most

$$\begin{aligned} (A_\rho + o(1))2^{p_k} \left(\Lambda_0 \frac{W_{k-1}}{w_0} + o(1) \right) \frac{2^{n_{k-1}}}{n_{k-1}} \\ = \left(\frac{A_\rho \Lambda_0}{w_0} + o(1) \right) \frac{2^n}{n}. \end{aligned}$$

The resulting g depends only on f , not on the requested input tuple, and the construction for T_k copies handles every smaller t by fixing unused inputs. Since $\Lambda_0 = 1$, letting $\alpha \downarrow 0$ proves (39). $\square$

Because the maximum defining $\Lambda_\gamma^{(\rho)}$ includes $t = 1$, (34) also gives

$$\Lambda_\gamma^{(\rho)} \geq 1 - H_2(\rho).$$

We therefore conjecture the approximate analogue of (28):

$$\boxed{\Lambda_\gamma^{(\rho)} = 1 - H_2(\rho) \quad \text{for every fixed } 0 \leq \rho < 1/2, 0 \leq \gamma < 1.}$$

As in the exact case, the missing factor is not a packing loss: it is the suffix-denominator loss caused by synthesizing the induced resource functions separately. The entropy gain also uses the freedom to choose the bad set after seeing f ; approximation on a prescribed set of locations is a different partial-function problem.

C Circuit implementation details

This appendix records the gate-level details behind the scheduler and routing lemmas. The main proof uses only their statements and costs.

C.1 Projective-direction indexing

Proof of Lemma 5.1. Normalize a nonzero vector so that its first nonzero coordinate is one. The normalized directions whose first nonzero coordinate is in position $h \in \{1, \dots, \ell\}$ have the form

$$(0, \dots, 0, 1, w_{h+1}, \dots, w_\ell)$$

and form a block of size $q^{\ell-h}$. Let

$$B_h = \sum_{r=1}^{h-1} q^{\ell-r}, \quad \text{val}_q(w_{h+1}, \dots, w_\ell) = \sum_{r=h+1}^{\ell} \text{int}(w_r) q^{\ell-r},$$

where $\text{int}(w_r) \in [q]_0$ is the canonical binary representation. The rank is

$$B_h + \text{val}_q(w_{h+1}, \dots, w_\ell).$$

The blocks are consecutive and their total size is D_q .

To rank, find the first nonzero coordinate, invert it, normalize the vector, and concatenate the b -bit tail coordinates. To unrank, compare the rank with the ℓ hardwired block intervals, subtract the appropriate B_h , and split the remainder into base- q digits. Since $q = 2^b$, the last step is bit slicing. The field operations and comparisons from the circuit conventions give size $\text{poly}_\ell(b)$ in both directions. $\square$

C.2 The least-missing-direction circuit

Proof of Lemma 5.2. At step j , fewer than D_q directions are forbidden, so a valid direction exists. This proves the combinatorial statement.

For the circuit bound, represent a projective direction by normalizing its first nonzero coordinate to one and use the ranking from Lemma 5.1. At stage j , for each $y \in U$, output the sentinel rank D_q when $y = u_j$ and otherwise output the projective rank of $y - u_j$. At the gate level, first multiplex a fixed nonzero vector in place of $y - u_j$ when $y = u_j$, rank that always-nonzero vector, and then overwrite the result by D_q on the equality branch. Use $\lceil \log_2(D_q + 1) \rceil$ bits per rank. There are $(j - 1)(q - 1) = O(jq)$ such entries.

Sort these ranks as $r_1 \leq \dots \leq r_s$. The least missing rank can be found without a dense D_q -bit table. Rank zero is a candidate if $s = 0$ or $r_1 > 0$; after position i , rank $r_i + 1$ is a candidate if

$$r_i < D_q - 1 \quad \text{and} \quad (i = s \text{ or } r_{i+1} > r_i + 1).$$

These tests are applied to the full sorted list, including duplicates and sentinels. Only the last occurrence of a non-sentinel rank can generate its successor, and a sentinel cannot generate a candidate.

A left-to-right prefix-OR circuit identifies the first candidate in this order. Multiplexers select its rank using $O(s \log D_q)$ gates.

Since $|U \setminus \{u_j\}| < D_q$, fewer than D_q distinct non-sentinel ranks are forbidden, so a candidate exists. Sorting, comparisons, and priority selection therefore cost $\tilde{O}_\ell(jq)$ gates. Unrank the selected direction using Lemma 5.1 and enumerate its $q - 1$ non-target line points, at cost $q \text{poly}_\ell(b)$. Unrolling the stages and summing over $j \leq g$ gives $\tilde{O}_\ell(g^2q)$. $\square$

C.3 Four-pass fixed-wire routing

Proof of Lemma 5.3. For scatter, create one slot record for each of the Cq^ℓ keys and combine them with the fewer than tq incidence records. Preserve a copy of every incidence record by free fan-out. Uniqueness inside a group puts at most one incidence at each key. After the first pass in the table, a slot therefore copies the preceding suffix exactly when a type-and-key equality test succeeds; otherwise it receives a fixed dummy suffix. The predecessor of the first position is defined to be a dummy record, so the boundary case uses the same guard. The second pass places the filled slots in canonical (group, y) order. Fixed wiring transposes this order to (y, group) , the input order used by the resource circuits.

For gather, combine the preserved incidences with one value record per key. The third pass places that value immediately before every matching incidence; a guarded local multiplexer copies it into the incidence. The fourth pass puts the first $t(q - 1)$ positions in the fixed request-and-line order required by the decoders. Thus absent incidences, the first array position, and all type boundaries are handled by fixed dummy data and the same guarded update, not by data-dependent wiring.

This is a constant number of sorting networks on $O(tq + Cq^\ell)$ records. Each record has width $\text{poly}_\ell(n, \log q, \log t, \log C)$, so the standard primitives from the circuit conventions give the stated bound. $\square$

References

- [1] Andreas Albrecht. On simultaneous realizations of Boolean functions, with applications. In *Parcella '88: Proceedings of the Fourth International Workshop on Parallel Processing by Cellular Automata and Arrays*, LNCS 342, pages 51–56. Springer, 1989. doi:10.1007/3-540-50647-0_102.
- [2] Alexander E. Andreev, Andrea E. F. Clementi, and José D. P. Rolim. Optimal bounds for the approximation of Boolean functions and some applications. *Theoretical Computer Science*, 180(1–2):243–268, 1997. doi:10.1016/S0304-3975(96)00217-4.
- [3] Kenneth E. Batcher. Sorting networks and their applications. In *Proceedings of the 1968 Spring Joint Computer Conference*, AFIPS Conference Proceedings 32, pages 307–314, 1968. doi:10.1145/1468075.1468121.
- [4] Arman Fazeli, Alexander Vardy, and Eitan Yaakobi. PIR with low storage overhead: Coding instead of replication. arXiv:1505.06241, 2015. <https://arxiv.org/abs/1505.06241>.
- [5] Gudmund Skovbjerg Frandsen and Peter Bro Miltersen. Reviewing bounds on the circuit size of the hardest functions. *Information Processing Letters*, 95(2):354–357, 2005. doi:10.1016/j.ipl.2005.03.009.
- [6] Lukas Holzbaur, Rina Polyanskaya, Nikita Polyanskii, and Ilya Vorobyev. Lifted Reed–Solomon codes with application to batch codes. In *2020 IEEE International Symposium on Information Theory*, pages 634–639. IEEE, 2020. doi:10.1109/ISIT44484.2020.9174402.
- [7] Yuval Ishai, Eyal Kushilevitz, Rafail Ostrovsky, and Amit Sahai. Batch codes and their applications. In *Proceedings of the 36th Annual ACM Symposium on Theory of Computing*, pages 262–271. ACM, 2004. doi:10.1145/1007352.1007396.
- [8] William Kretschmer. Quantum mass production theorems. In *18th Conference on the Theory of Quantum Computation, Communication and Cryptography*, LIPIcs 266, Article 10, pages 10:1–10:11, 2023. doi:10.4230/LIPIcs.TQC.2023.10.
- [9] Oleg B. Lupanov. On a method of circuit synthesis. *Izvestiya VUZ, Radiofizika*, 1(1):120–140, 1958. In Russian. <https://radiophysics.unn.ru/issues/1958/1/120>.
- [10] Wolfgang J. Paul. Realizing Boolean functions on disjoint sets of variables. *Theoretical Computer Science*, 2(3):383–396, 1976. doi:10.1016/0304-3975(76)90089-X.
- [11] Nicholas Pippenger. Information theory and the complexity of Boolean functions. *Mathematical Systems Theory*, 10:129–167, 1977. doi:10.1007/BF01683269.
- [12] Claude E. Shannon. The synthesis of two-terminal switching circuits. *Bell System Technical Journal*, 28(1):59–98, 1949. doi:10.1002/j.1538-7305.1949.tb03624.x.
- [13] Dietmar Uhlig. On the synthesis of self-correcting schemes from functional elements with a small number of reliable elements. *Mathematical Notes of the Academy of Sciences of the USSR*, 15(6):558–562, 1974. doi:10.1007/BF01152835.
- [14] Dietmar Uhlig. Networks computing Boolean functions for multiple input values. In M. S. Paterson, editor, *Boolean Function Complexity*, London Mathematical Society Lecture Note Series 169, pages 165–173. Cambridge University Press, 1992. doi:10.1017/CBO9780511526633.013.
- [15] Ingo Wegener. *The Complexity of Boolean Functions*. Wiley–Teubner, 1987. https://eccc.weizmann.ac.il/static/books/The_Complexity_of_Boolean_Functions/.